PARA Memory Pattern for nowu Agents

Inspired by: Paperclip's file-based PARA memory system (Projects, Areas, Resources, Archives)
Applies to: Agents that need persistent memory beyond a single session

---

When to Use This Pattern

Use PARA memory when an agent needs to remember things across heartbeat cycles that:

- Are not structural enough to be a `know` atom (too informal / too transient)
- Are not captured in artifacts (no explicit file home)
- Would lose context if the agent restarts (Memento Man problem)

PARA Structure for nowu

```
soul/memory/
├── PARA-projects.md      # Active focus: current epics, in-progress tasks
├── PARA-areas.md         # Ongoing responsibilities: architecture principles,
├── PARA-resources.md     # Reference material: links, patterns, prior resea
└── PARA-archive.md       # Completed cycles, closed epics, resolved decisio
```

PARA-projects.md Template

```
# Projects (Active Focus)
Last updated: YYYY-MM-DD

## Current Epic: epic-NNN — <title>
```

- Active intake: intake-NNN (status: READY_*FOR*_S1)
- Active task: task-NNN (status: IN_*PROGRESS)*
- *Story: story-NNN*
- *Estimated remaining: <Small / Medium / Large>*
- *Blocking issues: <none | describe>*
- *Last heartbeat progress: <one sentence>*

PARA-areas.md Template

 # Areas (Standing Responsibilities)
Last updated: YYYY-MM-DD

Architecture Principles
- <ADR-NNN: key constraint>
- <ADR-NNN: key constraint>

Known Agent Behavioral Patterns
- <agent-name>: tends to over-scope at S5, apply scope hammer
- <pattern>: <mitigation>

Human Preferences (Tacit Knowledge)
- Prefers short AC lists per task (max 4)
- Approves decisions in batches, not one-by-one

PARA-resources.md Template

 # Resources (Reference Material)
Last updated: YYYY-MM-DD

Patterns
- [Heartbeat pattern](ARTIFACT-heartbeat-skill.md)
- [Goal ancestry](ARTIFACT-goal-ancestry.md)

Prior Research
- <topic>: <one-line summary> (see disc-NNN-research.md)

```
## External References
- <URL or file>: <reason it matters>
```

Compounding Rule

At the end of each S9 capture cycle, the curator agent SHOULD:

1. Move completed tasks from `PARA-projects.md` → `PARA-archive.md`
2. Extract any durable behavioral patterns into `PARA-areas.md`
3. Create `know` atoms for any facts that are important enough to be cross-project knowledge

This prevents `PARA-projects.md` from becoming a graveyard. Only active items live in Projects.

Relationship to `know`

PARA memory is **informal, session-scoped working memory**. `know` atoms are **formal, structured, searchable knowledge**.

Write to PARA first. Graduate to `know` when a fact is:

- Referenced across more than one project
- Epistemically graded at INFORMED_ESTIMATE or above
- Worth tracking with version history

Why Not Just Use `know` for Everything?

`know` requires a `KnowledgeType`, `epistemic_grade`, and justification. For transient working notes (what happened in the last session, what the human prefers today), that overhead is wrong. PARA is the "scratch pad"; `know` is the "library".